

NLSR: Named-data Link State Routing Protocol

A K M Mahmudul Hoque
ahoque1@memphis.edu
University of Memphis

Syed Obaid Amin
soamin@memphis.edu
University of Memphis

Adam Alyyan
aalyyan@memphis.edu
University of Memphis

Beichuan Zhang
bzhang@cs.arizona.edu
University of Arizona

Lixia Zhang
lixia@cs.ucla.edu
University of California, Los Angeles

Lan Wang
lanwang@memphis.edu
University of Memphis

ABSTRACT

This paper presents the design of the Named-data Link State Routing protocol (NLSR), a routing protocol for Named Data Networking (NDN). Since NDN uses names to identify and retrieve data, NLSR propagates reachability to name prefixes instead of IP prefixes. Moreover, NLSR differs from IP-based link-state routing protocols in two fundamental ways. First, NLSR uses Interest/Data packets to disseminate routing updates, directly benefiting from NDN's data authenticity. Second, NLSR produces a list of ranked forwarding options for each name prefix to facilitate NDN's adaptive forwarding strategies. In this paper we discuss NLSR's main design choices on (1) a hierarchical naming scheme for routers, keys, and routing updates, (2) a hierarchical trust model for routing within a single administrative domain, (3) a hop-by-hop synchronization protocol to replace the traditional network-wide flooding for routing update dissemination, and (4) a simple way to rank multiple forwarding options. Compared with IP-based link state routing, NLSR offers more efficient update dissemination, built-in update authentication, and native support of multipath forwarding.

Categories and Subject Descriptors

C.2.2 [COMPUTER-COMMUNICATION NETWORKS]: Network Protocols—Routing Protocols

General Terms

Design, Security

Keywords

Routing, NDN, Trust Model

1. INTRODUCTION

The Named Data Networking (NDN) [3, 4] architecture is a fundamental paradigm shift from the current IP-based Internet architecture. Instead of carrying the destination IP address in each

packet, NDN puts a data name in each packet; a data consumer sends out an *Interest* packet whose name identifies the desired data, and the response is a *Data* packet containing the name, the data, and a signature by the original data producer. By explicitly naming and signing data, NDN enables features such as in-network caching, multipath forwarding, multicast data delivery and data authenticity.

For NDN to work well over wide-area networks, a routing protocol is needed to compute and install proper forwarding entries into an NDN node's forwarding table (FIB). Each FIB entry contains a name prefix and one or multiple next-hops, and is used to forward Interest packets whose names match the name prefix of the entry. While IP has to either use a single best next-hop or limit its forwarding to multiple equal-cost paths in order to avoid forwarding loops, NDN can utilize multiple paths freely because it has built-in loop prevention in the forwarding process. Thus an NDN network needs a routing protocol that can support name-based multipath routing.

We previously developed OSPFN [8], an extension to OSPF (Open Shortest Path First) for routing in NDN, and deployed it on the NDN testbed. OSPFN defines a new type of opaque link state announcement (LSA) to carry name prefixes in routing messages. It installs the best next-hop to each name prefix in the FIB; operators may manually configure a list of alternative next-hops for OSPFN to install in the FIB in addition to the best one. Although OSPFN can build a FIB with name prefixes, it has significant limitations. As in conventional IP-based routing protocols, OSPFN still uses IP addresses as router IDs, relies on GRE tunnels to cross legacy networks, and computes only a single best next-hop for each name prefix. Our experience from OSPFN deployment shows that managing IP addresses and tunnels are major operational problems, and the inadequate multipath support limits NDN's effectiveness.

In this paper we present the design of Named-data Link State Routing protocol (NLSR), which runs on top of NDN. In other words, NLSR uses NDN's Interest/Data packets to exchange routing messages. It is a link-state protocol as OSPF – link state advertisements are propagated throughout the entire network and each router builds a complete network topology. However, the route computation no longer produces just a single shortest path. It now *ranks* all policy-compliant next-hops and installs them into the FIB in order, essentially providing a name-based multipath routing table as input to NDN's forwarding strategy [9].

NLSR uses names instead of IP addresses to identify routers and links, therefore it can use any underlay communication channels (e.g., Ethernet, IP tunnels, TCP/UDP tunnels) for routing message exchanges. NLSR directly benefits from NDN's built-in data authenticity: since a routing update is carried in an NDN data packet and every NDN data packet carries a signature, a router can verify the signature of each routing message to ensure that it was gen-

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

ICN'13, August 12, 2013, Hong Kong, China.

Copyright is held by the owner/author(s). Publication rights licensed to ACM.

ACM 978-1-4503-2179-2/13/08 ...\$15.00.

erated by the claimed origin router and was not tampered during dissemination.

As the first distributed routing protocol on NDN, NLSR's design needs to answer the following important questions that are unique to applications running over NDN:

- **Naming:** how to name routers, links and routing updates.
- **Trust:** how to distribute routers' cryptographic keys and how to derive trust in these keys.
- **Information Dissemination:** how to disseminate routing updates in the network. While IP-based routing protocols *push* updates to other routers, NLSR routers need to *pull* the updates.
- **Multipath:** how to produce and rank the next-hops to facilitate multipath forwarding.

In this paper we describe our design choices and articulate the rationales behind these choices. Our goal is not to invent a new routing scheme as NLSR is essentially another link-state protocol, but rather, to demonstrate the feasibility and benefits of building a routing protocol using NDN.

Since NDN's adaptive, multipath forwarding is able to handle many packet delivery problems at the forwarding plane, the requirements on the routing plane is relaxed. For example, ensuring loop-freedom at the routing plane is no longer critical. Thus routing protocol designs that previously do not work in IP networks may now work in NDN networks, and new types of routing designs may also become possible. Exploring other types of routing designs would be our future work.

The remainder of the paper is organized as follows. The next section briefly introduces a few basic functions provided in an NDN network. Section 3 presents the design details. Section 4 provides the evaluation results. Section 5 discusses related work, and finally Section 6 concludes our work.

2. NDN PRIMER

NDN has three main components in its forwarding plane [3]: (1) *Forwarding Information Base (FIB)*: it stores the forwarding entries that direct Interest packets toward potential source(s) of matching Data. Unlike IP, it allows for a list of outgoing faces (next-hops) rather than a single one for each name prefix. The FIB is populated manually and/or by an NDN routing protocol in the control plane; (2) *Pending Interest Table (PIT)*: it stores the unsatisfied Interest packets and the faces on which they were received, so that Data packets can be routed back to the nodes interested in the data; and (3) *Content Store (CS)*: it is used for caching data.

When an Interest arrives at an NDN router, the CS is checked first for any matching data. If the Interest can be fulfilled by the CS, a Data packet is sent back on the face on which the Interest was received. Otherwise, it is added to the PIT. If there exists an entry with the same name, the new face number is added to the face list, so that a copy of the matching Data packet can be sent on all faces from which the Interest arrived. Finally, if the Interest does not have a matching PIT entry, the Interest is forwarded to the next-hop(s) based on the FIB. If multiple next-hops exist in a FIB entry, a module called "forwarding strategy" determines how the multiple routes will be used in forwarding Interests.

3. DESIGN

As a link-state protocol, NLSR disseminates Link State Advertisements (LSAs) to both build a network topology and distribute name prefix reachability. An NLSR router establishes and maintains adjacency relations with neighbor routers. Whenever it detects the failure or recovery of any of its links or neighbor pro-

cesses, it disseminates a new LSA to the entire network. Moreover, it advertises name prefixes from both static configuration and dynamic registration by local content producers. Whenever any name prefix is added or deleted, it also disseminates a new LSA. The latest version of the LSAs are stored in a Link State Database (LSDB) at each node.

Such topology and reachability dissemination may at first appear to be straight-forward as similar functions have been implemented in IP routing protocols. However, because we implement NLSR using NDN Interest and Data packets, the design must shift away from the familiar concepts of IP addresses and IP data pushing (i.e., any node can simply send any packet to any other node). Instead, we have to think in terms of data names and data retrieval. More specifically, we need a systematic naming scheme for routers and routing updates (Section 3.1). We also need to retrieve routing updates promptly without a priori knowledge of when an update may be generated, since a topology or name prefix change can happen any time (Section 3.3).

In terms of routing functionality, NLSR distinguishes itself from all previous link-state routing protocols in two aspects: (a) providing multiple routes to each name prefix, instead of a single shortest path; and (b) signing and verification of all LSAs to ensure that each router can originate only its own prefix and connectivity information. We present our route calculation algorithm in Section 3.4 and our trust model in Section 3.6.

As a preliminary step in developing NDN-based routing protocols, our initial design of NLSR is in the context of a single routing domain with a single authority that our trust model is built upon. We are in the process of deploying NLSR on the NDN testbed. We believe that this initial design and deployment experience can offer us a concrete stepping stone toward developing an NDN-based inter-domain routing protocol that incorporates routing policies and an inter-domain trust model.

3.1 Naming

Perhaps the most important piece in our design is a proper naming scheme for each element in the routing system and its corresponding public key. Based on the current network structures and operational practices, a hierarchical naming scheme seems best to capture the relationship among various components in the system, thus making it easy to identify routers belonging to the same network, as well as messages generated by a given routing process. It also helps associating keys with their corresponding owners.

In our design, each router is named according to the network it resides in, the specific site it belongs to, as well as an assigned router name, i.e., `/<network>/<site>/<router>`. For example, an ATT router in a PoP (point of presence) in Atlanta may be named `/ATT/AtlantaPoP1/router3`. This way, we know that if two routers share the same name prefix `/<network>`, they belong to the same network; and if they share the same prefix `/<network>/<site>`, they belong to the same site. This naming scheme makes it easy to filter out erroneous routing messages.

The NLSR process on a router is named after the router name: the router name is used as its prefix, followed by the process name NLSR, i.e., `/<network>/<site>/<router>/NLSR`. This NLSR routing demon name is used in periodic *info* messages between adjacent NLSR routers to detect the failure of either links or routing processes themselves (see Section 3.5).

Ideally, any routing updates originated by an NLSR process should have the process name as its prefix to easily identify the messages originator. In other words, the name for an LSA should begin with `/<network>/<site>/<router>/NLSR/LSA` to indicate that it is generated by the NLSR process. However, because

Table 1: Contents of an LSA

Type	Content
Adjacency LSA	# Active Links (N), Neighbor 1 Name, Link 1 Cost, ..., Neighbor N Name, Link N Cost
Prefix LSA	isValid, Name Prefix

our implementation uses CCNx Sync [5] and Repo [5] to disseminate LSA data, and CCNx repo imposes a constraint that all the data to be synchronized must share a common name prefix, our current implementation is confined to using a common prefix for the LSAs generated by all the routers. We name each LSA using the common prefix `/<network>/NLSR/LSA` (we call this `<LSA-prefix>`), and append `/<site>/<router>` to this prefix to differentiate LSAs originated by different NLSR routers.

3.2 LSAs

NLSR is designed to propagate two types of LSAs – Adjacency LSA and Prefix LSA. The Adjacency LSA is used to advertise all active links connecting one NDN router to its neighbors. The Prefix LSA, on the other hand, is used to advertise a name prefix that has been registered with the router. Their contents are shown in Table 1.

An Adjacency LSA has the name format `/<LSA-prefix>/<site>/<router>/LsType.1/<version>`, where `<router>` is the name of the router that originates the LSA and `<version>` indicates the ordering in the various versions of a particular LSA as it changes over time. It is currently implemented as the LSA origination times in microseconds from epoch time. However, similar to OSPF, sequence numbers can also be used for this purpose. As shown in Table 1, the Adjacency LSA contains all the *active* links of a router, each associated with a neighboring router’s name and a link cost. It is created at router startup time and whenever there is any status change in a router’s links, as detected by periodical “info” Interest messages (Section 3.5).

A Prefix LSA has the name format `/<LSA-prefix>/<site>/<router>/LsType.2/LsId.<ID>/<version>`. Note that each Prefix LSA advertises one name prefix. Since one router may have multiple name prefixes registered with it, it needs to announce multiple Prefix LSAs, using a unique LSA ID¹ in their name to differentiate them. The rationale for this design decision is that bundling all the name prefixes of a router in a single LSA may make it too large to be transported in one message and also inefficient to update (even if only one prefix is added or removed, all the other prefixes in the same LSA need to be advertised again). Each Prefix LSA contains a flag *isValid* (set to 1 initially) and the name prefix to be advertised (Table 1). When a name prefix is de-registered, NLSR updates the corresponding prefix LSA by setting *isValid* to 0, and disseminates the new LSA to other nodes. An NLSR node receiving this LSA will delete this name prefix from its LSDB and update its FIB accordingly.

In order to remove obsolete LSAs caused by router crashes, every router periodically refreshes each of its advertised LSAs by generating a newer version. Every LSA has a lifetime associated with it, and will be removed from the LSDB when the lifetime expires. Therefore if a router crashes, its LSAs will not persist in other routers’ LSDBs. Note that route calculation should not be impacted by the obsolete LSAs in NLSR – if a router crashes, its neighbors will update the status of their LSAs so traffic will not be directed over those links. Since we do not use the refreshes to handle packet losses or state corruption (CCNx Sync handles it) and the obsolete

¹The LSA ID can be manually configured or calculated based on the name prefix (e.g., a hash of the name prefix).

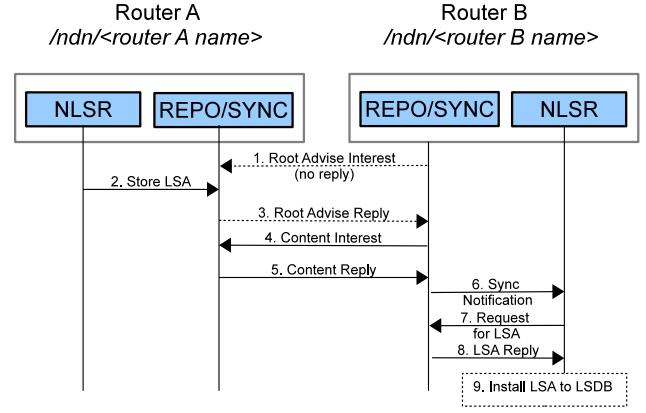


Figure 1: LSA dissemination from router to router via CCNx Sync/repo (dotted line represents periodic messages.)

LSAs do not affect route calculation, these refreshes should be sent at a relatively long interval, e.g., on the order of days.

3.3 LSDB Synchronization

To simplify our design conceptually, we decided to view the LSDB as a collection of data, and the LSA dissemination problem as a data synchronization problem of the LSDBs maintained by the routers. Routers periodically exchange their hashes of the LSDB to detect inconsistencies and recover from them. This hop-by-hop synchronization approach avoids unnecessary flooding to the network – when the network is stable, only one hash, instead of all the LSAs, is exchanged between neighbors. Moreover, it is also receiver-driven, meaning that a router will request LSAs only when it has CPU cycles. Thus it is less likely a router will be overwhelmed by a flurry of updates.

Our current implementation uses the CCNx synchronization protocol, or *Sync* [5], to disseminate the LSAs to the neighboring routers. *Sync* is associated with the CCNx repository (or *Repo*). It allows applications to define collections of named data in a repo, called *slices*, which are kept in sync with identically defined slices in neighboring repos. *Sync* computes a hash tree over all the data in a slice and exchanges the root hash between neighbors to detect inconsistencies. If the hash values do not agree, two neighboring nodes then exchange the hash values of nodes on the next tree level until they detect the specific leaf nodes (data) causing the problems. They then exchange the data to reach consistency.

Figure 1 shows how an LSA is disseminated in the network. To synchronize the slice containing LSAs, the *Sync* protocol periodically sends special Interest messages, called *Root Advise*, with the combined hash value of the slice to the neighboring nodes (step 1). When Router A’s NLSR creates an LSA and writes it in the *Sync* slice (step 2), its hash value becomes different from that of Router B, which causes Router A’s *Sync* to reply to the *Root Advise* Interest from Router B with the new hash value of its local slice (step 3). Router B’s *Sync* then compares the hashes and recursively requests for the next level hashes that cause the differences. Eventually, Router B’s *Sync* identifies the data that needs to be synchronized (LSAs in the context of NLSR) and retrieves them using Interest messages (step 4 and 5). The *Sync* on Router B then sends the data *name* to the local NLSR agent (step 6), which fetches the data from the local repo (step 7 and 8) and updates its LSDB (step 9).

Each *Root Advise* Interest has a lifetime and a new *Root Advise* is sent when the lifetime expires. Such periodic transmission is designed to handle the loss of *Root Advise* Interests, and thus reduce